

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 import xgboost as xgb
7 from sklearn.metrics import mean_squared_error
8 color_pal = sns.color_palette()
9 plt.style.use('fivethirtyeight')
```

```
1 df = pd.read_csv('/content/AEP_hourly.csv')
2 df = df.set_index('Datetime')
3 df.index = pd.to_datetime(df.index)
```

```
1 df.head()
2
```

1 to 5 of 5 entries Filter ?


Datetime	AEP_MW	hour	dayofweek	quarter	month	year	dayofyear
2004-12-31 01:00:00	13478.0	1	4	4	12	2004	366
2004-12-31 02:00:00	12865.0	2	4	4	12	2004	366
2004-12-31 03:00:00	12577.0	3	4	4	12	2004	366
2004-12-31 04:00:00	12517.0	4	4	4	12	2004	366
2004-12-31 05:00:00	12670.0	5	4	4	12	2004	366

Show 25 per page




Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.


Distributions




2-d distributions



Time series



Values



```
1 df.tail()
```



1 to 5 of 5 entries Filter 📄 ?


Datetime	AEP_MW
2018-01-01 20:00:00	21089.0
2018-01-01 21:00:00	20999.0
2018-01-01 22:00:00	20820.0
2018-01-01 23:00:00	20415.0
2018-01-02 00:00:00	19993.0

Show 25 per page




Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Distributions



Values



```
1 df.describe()
2
```



1 to 8 of 8 entries Filter 📄 ?


index	AEP_MW
count	121273.0
mean	15499.513716985644
std	2591.399065407914
min	9581.0
25%	13630.0
50%	15310.0
75%	17200.0
max	25695.0

Show 25 per page




Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.


Distributions



Values




```
1 df.plot(style='.',
2         figsize=(15, 5),
3         color=color_pal[2],
4         title='PJME Energy use in MW')
5 plt.show()
6
```



1 Start coding or generate with AI.


TRAIN/TEST SPLIT

```
1 train = df.loc[df.index < '01-01-2015']
2 test = df.loc[df.index >= '01-01-2015']
3
4 fig, ax = plt.subplots(figsize=(15, 5))
5 train.plot(ax=ax, label = 'Training set',title='Data Train/Test Split')
6 test.plot(ax=ax, label = 'Test Set')
7 ax.axvline('01-01-2015', color = 'black', ls='--')
8 ax.legend(['Training Set','Test Set'])
9 plt.show()
```



```
1 df.loc[(df.index > '01-01-2010') & (df.index < '01-08-2010') ].plot()
```

<Axes: xlabel='Datetime'>



FEATURE CREATION

```
1 df.index.hour
```

```
Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,
        ...
        15, 16, 17, 18, 19, 20, 21, 22, 23,  0],
      dtype='int32', name='Datetime', length=121273)
```

```
1 def create_features(df):
2
3     df['hour'] = df.index.hour
4     df['dayofweek'] = df.index.dayofweek
5     df['quarter'] = df.index.quarter
6     df['month'] = df.index.month
7     df['year'] = df.index.year
8     df['dayofyear'] = df.index.dayofyear
9     return df
10
11 df = create_features(df)
```

VISUALIZE OUR FEATURE /TARGET RELATIONSHIP

```
1 print(df.columns)
2
```

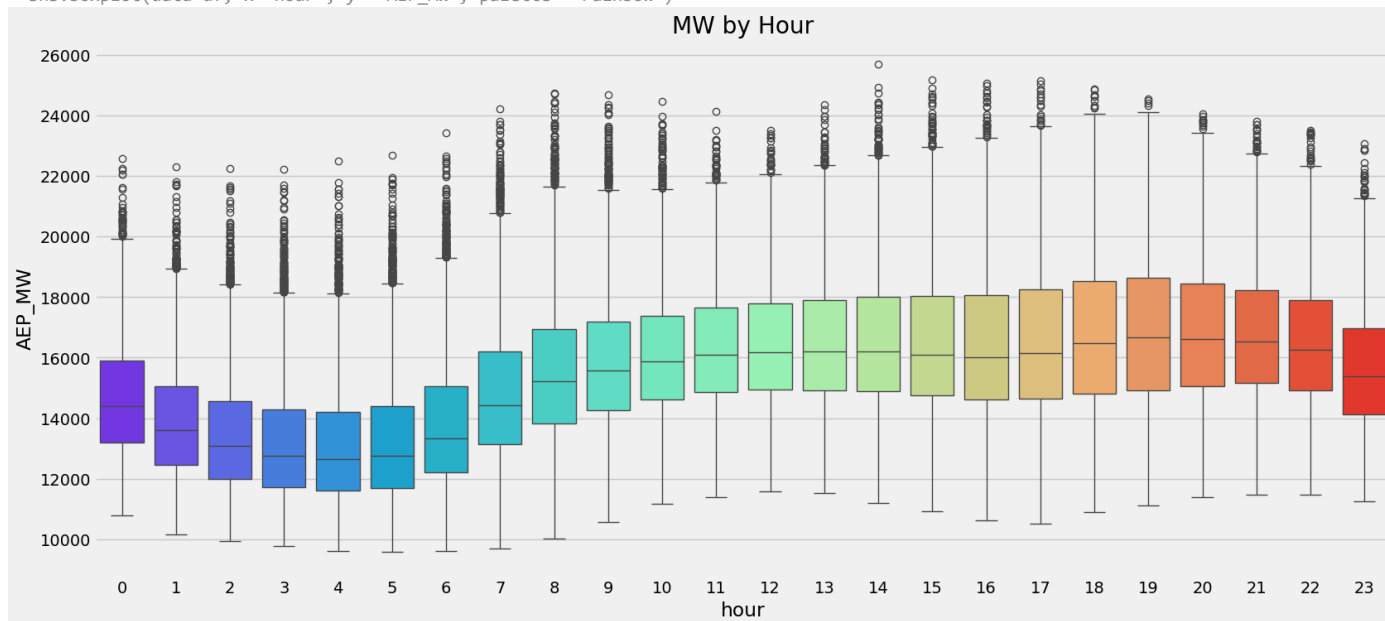
```
Index(['AEP_MW', 'hour', 'dayofweek', 'quarter', 'month', 'year', 'dayofyear'], dtype='object')
```

```
1 fig, ax = plt.subplots(figsize=(18,8))
2 sns.boxplot(data=df, x='hour', y='AEP_MW', palette='rainbow')
3 ax.set_title('MW by Hour')
4 plt.show()
```

```

<ipython-input-37-148aa888781f>:2: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`
sns.boxplot(data=df, x='hour', y='AEP_MW', palette='rainbow')

```



```


1 fig, ax = plt.subplots(figsize=(18,8))
2 sns.boxplot(data=df, x='month', y='AEP_MW', palette='Greens')
3 ax.set_title('MW by Month')
4 plt.show()

```

 <ipython-input-38-c92272204d98>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.boxplot(data=df, x='month', y='AEP_MW', palette='Greens')
```



✓ CREATE A MODEL

```
1 train = create_features(train)
2 test = create_features(test)
3 FEATURES = ['hour', 'dayofweek', 'quarter', 'month', 'year', 'dayofyear']
4 TARGET = 'AEP_MW'
5
6
```

 See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['quarter'] = df.index.quarter
<ipython-input-35-e87537c8a040>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['month'] = df.index.month
<ipython-input-35-e87537c8a040>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
<ipython-input-35-e87537c8a040>:4: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['dayofweek'] = df.index.dayofweek
```

```
<ipython-input-35-e87537c8a040>:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['quarter'] = df.index.quarter
```

```
<ipython-input-35-e87537c8a040>:6: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['month'] = df.index.month
```

```
<ipython-input-35-e87537c8a040>:7: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['year'] = df.index.year
```

```
<ipython-input-35-e87537c8a040>:8: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a

```
df['dayofyear'] = df.index.dayofyear
```

```
1 X_train = train[FEATURES]
```

```
2 y_train = train[TARGET]
```

```
3
```

```
4 X_test = test[FEATURES]
```

```
5 y_test = test[TARGET]
```

```
1 !pip install -U xgboost
```



Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (2.1.4)

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.26.4)

Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)

Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.13.1)

```
1 import xgboost
```

```
2 print(xgboost.__version__)
```



```
2.1.4
```

```
1 reg = xgb.XGBRegressor(n_estimators=1000,early_stopping_rounds=50, learning_rate=0.001)
```

```
2
```

```
3 reg.fit(X_train, y_train,
```

```
4     eval_set=[(X_train, y_train), (X_test, y_test)],
```

```
5     verbose=100
```

```
6 )
```

```
7
```

```
8
```

```

[0]    validation_0-rmse:2571.77747    validation_1-rmse:2683.12724
[100]  validation_0-rmse:2418.78537    validation_1-rmse:2543.72673
[200]  validation_0-rmse:2285.21737    validation_1-rmse:2424.24267
[300]  validation_0-rmse:2168.46028    validation_1-rmse:2320.84646
[400]  validation_0-rmse:2066.27063    validation_1-rmse:2231.99828
[500]  validation_0-rmse:1977.35291    validation_1-rmse:2156.44867
[600]  validation_0-rmse:1900.68213    validation_1-rmse:2095.13001
[700]  validation_0-rmse:1832.44385    validation_1-rmse:2040.87265
[800]  validation_0-rmse:1772.78374    validation_1-rmse:1993.65343
[900]  validation_0-rmse:1720.60363    validation_1-rmse:1949.47821
[999]  validation_0-rmse:1674.54889    validation_1-rmse:1911.14625

```

```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytreet=None, device=None, early_stopping_rounds=50,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.001, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=1000, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)

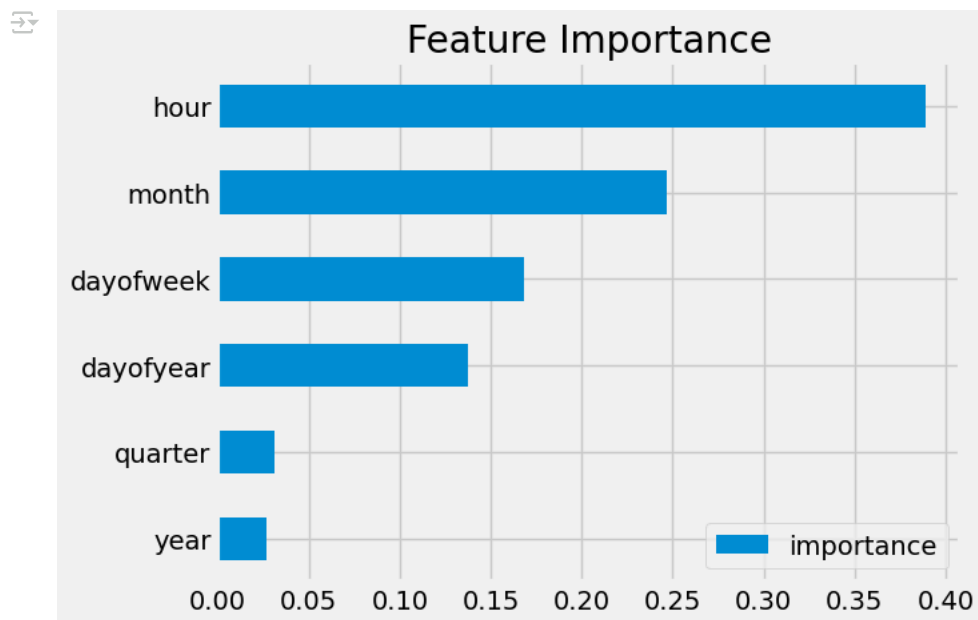
```

FEATURE IMPORTANCE

```

1 fi = pd.DataFrame(data=reg.feature_importances_,
2                   index=reg.feature_names_in_,
3                   columns=['importance'])
4 fi.sort_values('importance').plot(kind='barh', title='Feature Importance')
5 plt.show()

```



FORECAST ON TEST

```
1 test['prediction'] = reg.predict(X_test)
```

```

<ipython-input-62-33b5c0e814cb>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
test['prediction'] = reg.predict(X_test)
```



```
1 df.columns
```

```
Index(['AEP_MW', 'hour', 'dayofweek', 'quarter', 'month', 'year', 'dayofyear'], dtype='object')
```

```
1 print(type(test))
2 print(test.head())
3
```


```
<class 'pandas.core.frame.DataFrame'>
      AEP_MW  hour  dayofweek  quarter  month  year \
Datetime
2015-01-01 00:00:00 16375.0    0        3        1     1  2015
2015-12-31 01:00:00 12415.0    1        3        4    12  2015
2015-12-31 02:00:00 12087.0    2        3        4    12  2015
2015-12-31 03:00:00 12010.0    3        3        4    12  2015
2015-12-31 04:00:00 11972.0    4        3        4    12  2015
```

```
      dayofyear  prediction
Datetime
2015-01-01 00:00:00      1 15923.991211
2015-12-31 01:00:00    365 14933.623047
2015-12-31 02:00:00    365 14924.480469
2015-12-31 03:00:00    365 14924.480469
2015-12-31 04:00:00    365 14924.480469
```

```
1 print(f"X_test shape: {X_test.shape}")
2 print(f"test shape: {test.shape}")
3
```

```
X_test shape: (31440, 6)
test shape: (31440, 8)
```

```
1 df = df.merge(test[['prediction']], how='left', left_index=True, right_index=True)
2 ax = df[['AEP_MW']].plot(figsize=(15, 5))
3 df['prediction'].plot(ax=ax, style='.')
4 plt.legend(['truth data', 'predictions'])
5 ax.set_title('raw data and prediction')
6 plt.show()
```



```
1 score = nn.smf(mean_squared_error(test['AEP_MW'], test['prediction']))
Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.
```